

Multi-Course Accessibility Audit Protocol

Multi-Course Accessibility Audit Protocol

A survey method for auditing digital course shells against WCAG 2.1 Level AA

Version 1.1 · Saroj Lamichhane · Licensed CC BY 4.0

Why an error count is the wrong output

Automated checkers report failures. A large course shell returns hundreds; an institution returns tens of thousands. The number is accurate and close to useless, for three reasons.

It has no severity. A missing document title and an image-only PDF of the midterm both count as one. One is a formality, the other locks a student out of an assessment.

It has no owner. An error list arrives on an administrator's desk with no indication of who fixes what, so it goes to whoever is least able to refuse.

It has no sequence. Faculty asked to fix everything fix nothing, because there is no defensible place to start.

This protocol produces a different output: a course-by-course remediation sequence, ordered by consequence to students, with a record of who has to do the work. Its stated purpose is not a compliance score. It is to identify what faculty need to be taught.

Why not simply use the accessibility tool inside the LMS

Most learning management systems now ship with an accessibility checker, and many institutions license a scoring product such as Blackboard Ally or UDOIT. Use them. They are fast, and they catch things a human reviewer misses on a tired afternoon.

They are not an audit, for a reason the W3C states directly: automated tools cannot determine whether content conforms to the standard, and evaluation requires human judgment. The specific gaps matter.

They test presence, not quality. A checker records that alternative text exists. It cannot tell you the alt text reads "image1.png", or that on a worksheet asking students to identify a structure, the alt text names the answer. Presence is machine-checkable. Adequacy is not.

They score files, not courses. A course can score well and still be unusable: every document individually acceptable, no consistent structure, navigation that makes no sense. Students experience the course, not the file.

They cannot judge reading order or caption accuracy. A tagged PDF with a scrambled tag tree passes as tagged. An automatic caption track passes as captioned, including when it renders your discipline's vocabulary as nonsense. Both need someone to look.

They produce a number that gets optimised. When a course carries a visible score, effort migrates to whatever raises it fastest, which is rarely the blocking failure. A green score on a course whose midterm is an untagged scan is worse than no score, because it retires the question.

They do not tell you who fixes it. That is the field which determines whether anything happens, and no product supplies it.

Use the automated output as an input to Step 3 below, not as a substitute for it. Let the machine find candidates; let a person classify them.

What you need before starting

Reviewer access to the course shells you intend to audit, and a copy of the review workbook linked at the end of this page. The free WAVE browser extension, a PDF reader that shows document properties and tags, and a screen reader you are willing to spend five minutes with, NVDA on Windows or VoiceOver on macOS.

You do not need a commercial accessibility platform. If your institution has one, use its output as an input.

Budget roughly 45 to 90 minutes per course. This is a survey instrument, not a conformance evaluation, and the distinction matters: a full evaluation of a single course can take a day, which is why institutions never finish one.

Step 1. Sample rather than enumerate

Auditing every file in every course is how audits die. This instrument samples twelve files per course, chosen to span the content types students actually meet: PDFs, Word documents, slide decks, video, links, and the LMS pages, navigation and structure themselves.

Twelve is enough to characterise a course and small enough that the audit finishes. If eight of twelve files fail on heading structure, you do not need the other forty to know what the course needs.

Choose courses deliberately too. Sample across disciplines rather than within one, because failure patterns differ sharply by field: chemistry and mathematics fail on notation, nursing on scanned clinical handouts, business on embedded spreadsheets and decks, humanities on long unstructured reading PDFs. Weight toward required, first-year and gateway courses, and anything with a laboratory or clinical component.

Step 2. Record the same eight checks on every file

Consistency matters more than sophistication, because the value appears when you compare across courses and across semesters.

Check	What fails it
Heading and list structure	Visual formatting used in place of real

	heading styles and list tools
Descriptive links	“Click here”, bare URLs, link text that does not name the destination
Alternative text	Missing, or present but empty of meaning
Tables	No header row, merged or empty cells, layout tables
Video captions	Absent, or automatically generated and uncorrected
Document title	Not set in document properties
Tags	Not defined, the standard failure in PDFs
Reading order	Not set, or not matching the visual order

Alongside the checks, record the semester, course identifier, file type, file name, notes on what was found and what was done, an overall note on the course, the date, and the reviewer.

The free-text columns carry more weight than they look. “Slide deck embedded through an inline viewer, screen readers cannot reach the text” is the finding that changes institutional practice, and no checkbox captures it.

Step 3. Classify findings by consequence

Every finding goes into one of three tiers. This is the step that converts a list into a plan.

Blocking. A student using assistive technology cannot reach the content at all. Image-only PDFs with no text layer. Slide decks embedded through an inline viewer. Video with no captions and no transcript. Equations and chemical structures inserted as pictures. Anything unreachable by keyboard.

Degrading. Reachable, but substantially harder to use. No heading structure. Layout tables, or data tables with no header row. Link text that names nothing. Alternative text present but meaningless. Colour as the only signal.

Cosmetic. Real failures that do not prevent access. Contrast marginally below threshold. Missing document title or language.

Cosmetic findings still get fixed. They are simply not where a programme starts, and putting them in the same list as a locked assessment is what makes a report unusable.

Step 4. Build the priority matrix

1. Blocking failures on assessments and safety-critical material. A student who cannot read a quiz question or a laboratory safety step is harmed immediately, and institutional exposure is greatest.
2. Blocking failures on required readings and core content in high-enrollment required courses.
3. Blocking failures elsewhere.

4. Degrading failures, in the same order.
5. Cosmetic failures, batched and fixed institutionally where possible, since many are template problems affecting every course at once.

Two adjustments are worth making. Where a single template or theme causes a failure across many courses, promote it: fixing one master template can clear thousands of findings. And where a course is already scheduled for redesign, defer everything except blocking assessment failures, because remediating a course you are about to rebuild is wasted work.

Step 5. Report tiers, not percentages

Record baseline counts by tier, then recount after remediation using the same instrument. Report the change in blocking failures as the headline.

Resist the single compliance percentage. It invites the conclusion that ninety-four percent is nearly done, when the remaining six percent may be every assessment in the nursing programme.

Step 6. Convert findings into training

This is the objective the instrument was built to serve. An audit producing only a work order leaves the institution exactly as capable as it was.

Every recurring degrading failure is a curriculum item. The pattern is consistent: heading structure, alternative text quality, table headers, link text, and PDF export method account for the large majority of degrading findings, and each is teachable in under an hour. That is what the free faculty course is built from, and its module order follows the frequency of audit findings rather than the order of the success criteria.

Fix blocking failures centrally, because they are urgent and often technical. Teach degrading failures, because they recur every term and no central team can keep up with them.

Step 7. Schedule the re-check

Courses decay. A course remediated in spring acquires new uploads by autumn, and a course that changes instructors usually reverts. Set a review interval, re-sample twelve files, and compare against the previous record. The workbook is structured by semester for exactly this reason.

What this protocol is not

It is not a conformance evaluation and should not be described as one. It does not produce a WCAG 2.1 Level AA conformance claim. An institution needing a formal conformance report will need a full evaluation, for which the W3C's Website Accessibility Conformance Evaluation Methodology is the appropriate framework.

What this produces is triage: a defensible answer to “how bad is it, what do we fix first, and what do we need to teach,” arrived at quickly enough that the answer is still true when you act on it.

References

The approach below is standard practice rather than invention. The sampling model, the insistence on human judgment, and the treatment of automated output as an input all follow published guidance.

- [Website Accessibility Conformance Evaluation Methodology \(WCAG-EM\) 1.0](#), W3C. The reference methodology for structured evaluation, including the use of a selected sample rather than exhaustive review.
- [Selecting Web Accessibility Evaluation Tools](#), W3C WAI. States that automated tools cannot determine conformance and support rather than replace human evaluation.
- [Easy Checks, A First Review of Web Accessibility](#), W3C WAI. The preliminary-review model this protocol follows.
- [Involving Users in Evaluating Web Accessibility](#), W3C WAI. Why checker output and lived experience diverge.
- [The WebAIM Million](#), WebAIM. Annual large-scale evidence on the limits of automated detection.
- [ADA Title II web and mobile accessibility rule](#), U.S. Department of Justice. The obligation and the compliance dates.
- [Web Content Accessibility Guidelines \(WCAG\) 2.1](#), W3C. The standard itself.